

# Tweaking Neural Net

## Parameters



Global  
HD

MEMECENTER.COM



**RECENTLY RELEASED,  
UNPROVEN DEEP  
LEARNING MODEL**



**STANDARD, STRONG  
BASELINE MODEL**

# CNNs – General Layout

- One-time setup
  - **Architecture**
  - Activation functions (sigmoid, ReLU, ...)
  - Regularization (batch norm, dropout)
- Training
  - Data collection: Preprocessing, Augmentation
  - Loss function: cross entropy loss...
  - Training via SGD (update rules)

# The Humble Beginnings

## LeNet-5

State-of-the-art performance on MNIST digit recognition ( $< 1\%$ )





Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

[LeCun, Y., Bottou, L., Bengio, Y., & Haffner, P. \(1998\). Gradient-based learning applied to document recognition. Proceedings of the IEEE, 86\(11\), 2278-2324.](#)



Conv 1

MaxPool 1

Conv 2

MaxPool 2

Conv 3,4,5

MaxPool 3

FCN 6,7,8

96 11x11  
Stride 4  
Pad 0

3x3  
Stride 2

256 5x5  
Stride 2  
Pad 2

3x3  
Stride 2

384 3x3  
Stride 1  
Pad 1

384 3x3  
Stride 1  
Pad 1

256 3x3  
Stride 1  
Pad 1

3x3  
Stride 2

4096 4096 1000

55x55x96

27x27x96

27x27x256

13x13x256

13x13x384

13x13x384

13x13x256

6x6x256

4096 4096 1000

# VGG

## Q: Why use smaller filter?

Let's do the other way around: Why use larger filters?


□ Receptive field!

Stack of 3 3x3 stride 1 convolutional layers has **same effective receptive field** as 7x7 layer!

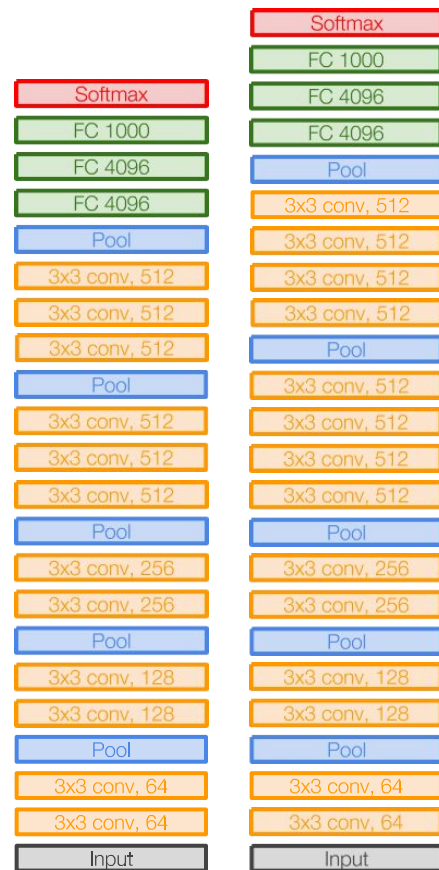
But deeper □ **More non-linearities**

**Fewer parameters!**

$3 \cdot (3^2 C)$  vs  $7^2 C$  for  $C$  channels per layer



AlexNet




VGG16

VGG19

Simonyan, K., & Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. arXiv preprint arXiv:1409.1556.

# ResNet

**Intuition:** Deeper network should perform at least as well as a shallower version.  
When implemented naively, does this intuition hold?

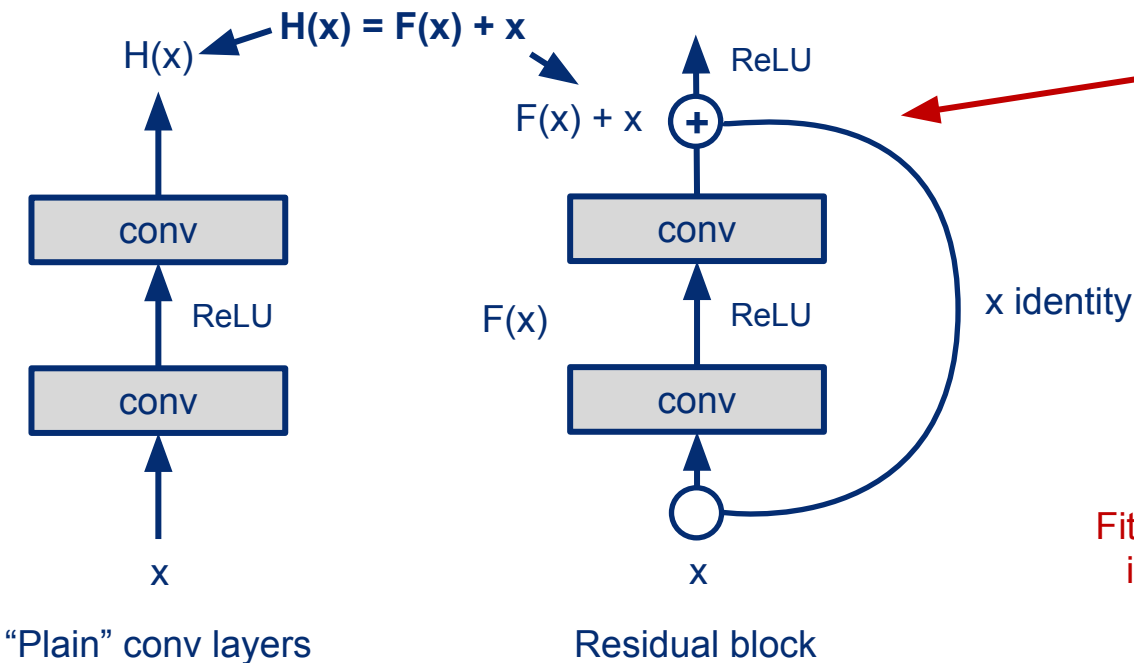


He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. *IEEE conference on computer vision and pattern recognition* (pp. 770-778).

# ResNet

## Solution

Use network layers to fit residual mapping (rather than desired mapping directly)



**Q: What about gradients at addition nodes?**

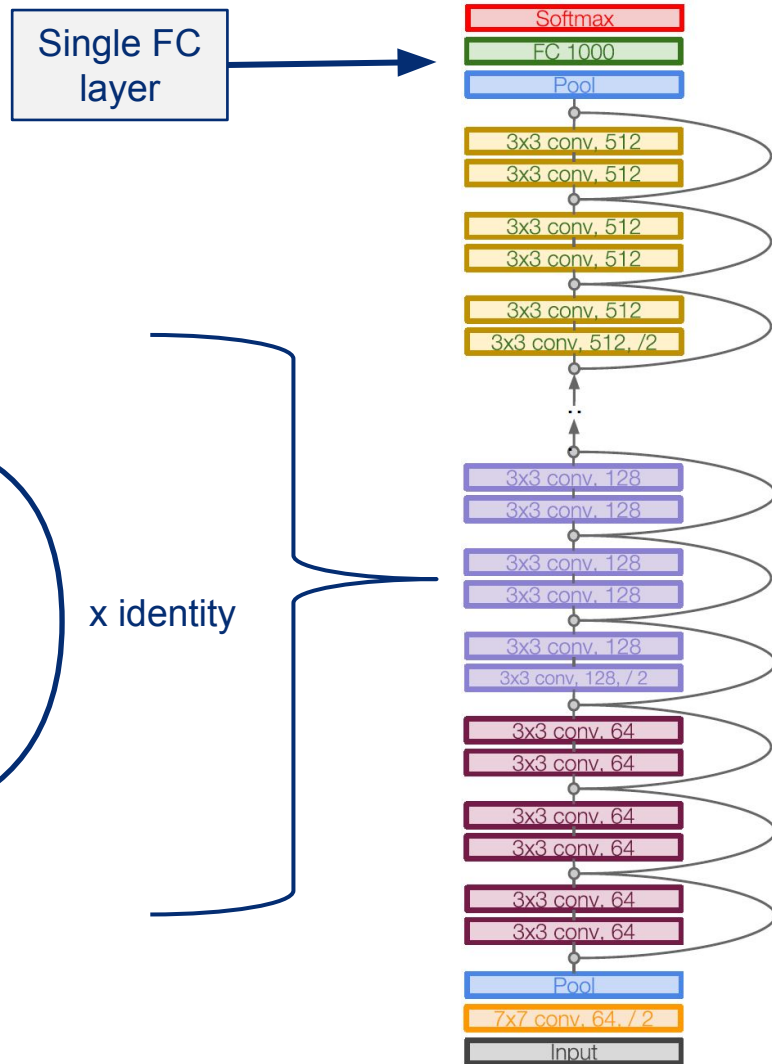

- Add: Distribute gradient!
  - Gradient flows unhindered
  - "Gradient highway"
- Makes training easier!

Fit residual  $F(x) = H(x) - x$  instead of  $H(x)$  directly

# ResNet

## Full architecture

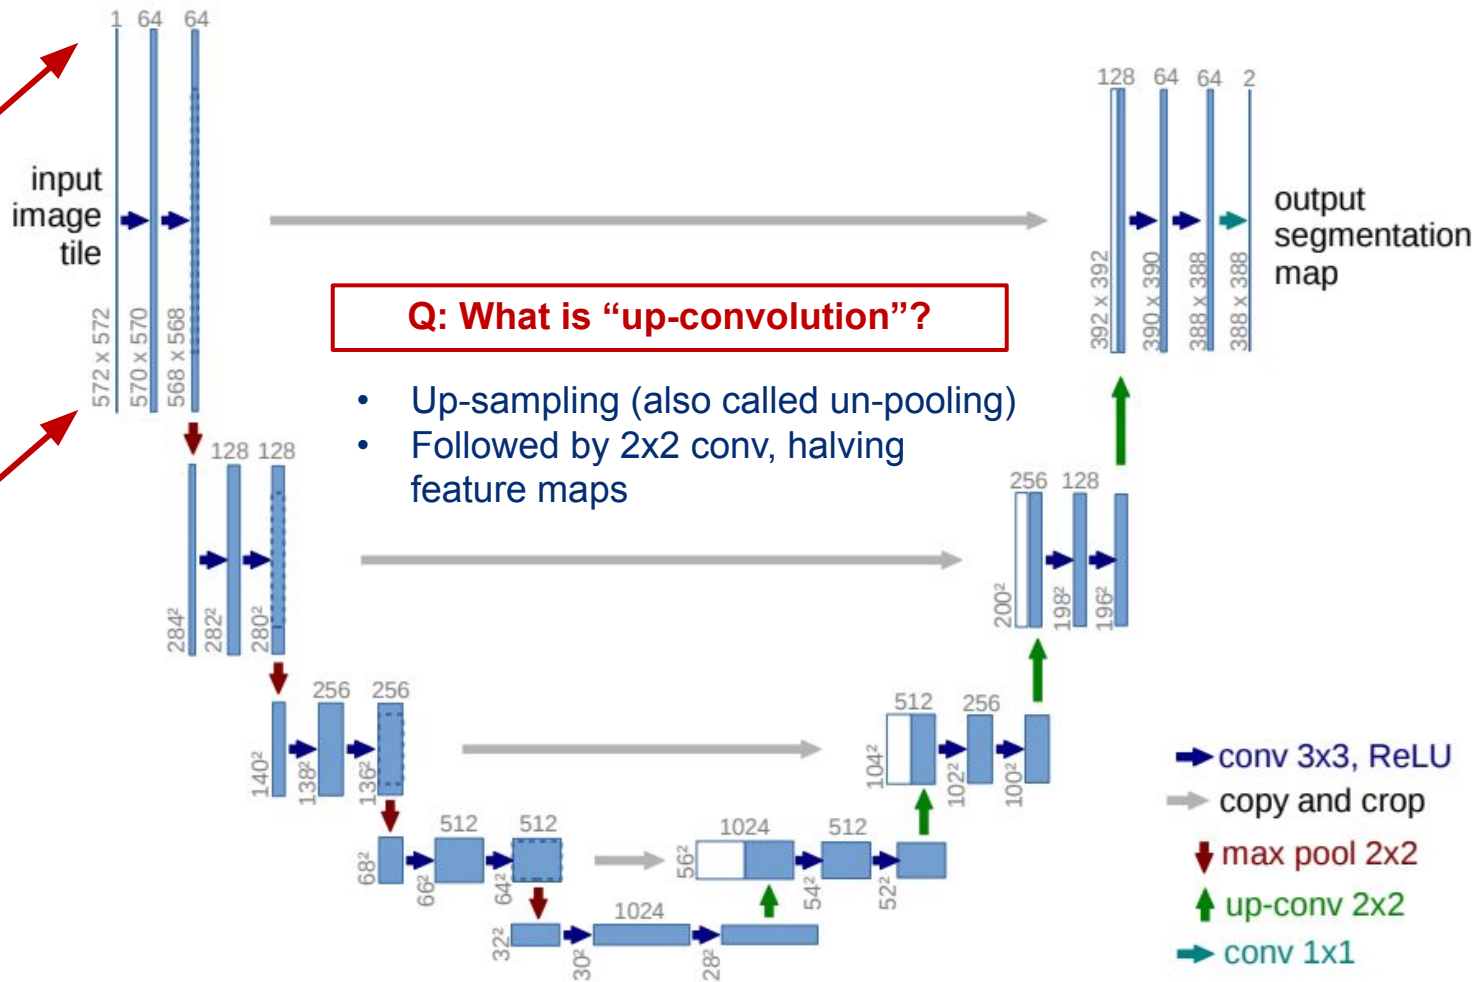
- Stack residual blocks
- Residual block has two 3x3 conv layers
- Periodically, double # filters and down-sample with stride 2
- Additional conv layer at beginning
- Only one FC layer





# feature maps

Size





# Upsampling

Goal: Increase the spatial dimensions of your matrix

Methods (of interpolation): Nearest neighbor,

## Nearest Neighbor

1	2
3	4



1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Input: 2 x 2

Output: 4 x 4

- **mode** (*str*, optional) – the upsampling algorithm: one of 'nearest', 'linear', 'bilinear', 'bicubic' and 'trilinear'. Default: 'nearest'



# Upsampling

Goal: Increase the spatial dimensions of your matrix

Methods (of interpolation): Nearest neighbor,

## Nearest Neighbor

1	2
3	4



1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

Input: 2 x 2

Output: 4 x 4

When would you want to use which?

- **mode** (*str*, optional) – the upsampling algorithm: one of 'nearest', 'linear', 'bilinear', 'bicubic' and 'trilinear'. Default: 'nearest'


# Upsampling

How many parameters does this have?

# Upsampling

How many parameters does this have? 0

There's nothing to learn... do we want to learn something?



# Transposed Convolution

Basically the same thing as upsampling but... with learnable parameters

There's nothing to learn... **do we want to learn something?**

