

ONE DOES NOT SIMPLY

SAMPLE THE LATENT SPACE




Autoencoding

Denoising Autoencoder

$$\theta = \arg \min_{\hat{\theta}} d[x, g(f(\tilde{x}, \hat{\theta}), \hat{\theta})]$$

- “Encoder”: $f()$
- “Decoder”: $g()$
- **Q: What is \tilde{x} ?**

→ Heavily corrupted versions of x ! Corruption: Zero-ing, noise, etc...





[Vincent, P., Larochelle, H., Bengio, Y., & Manzagol, P. A. \(2008, July\). Extracting and composing robust features with denoising autoencoders. ICML.](#)
[Vincent, P. et al. \(2010\). Stacked denoising autoencoders: Learning useful representations in a deep network with a local denoising criterion. JMLR.](#)

Learning Meaningful Feature Representations from Data

After training: Two sub-modules!

- Image generator
 - Remove encoder
 - Randomly generate latent vector z
 - Generate output sample
 - ??
 - Profit



Q: Why?

Embedding shows clusters

→ Makes sense: Distinct encoding for distinct image appearance

But: Embedding is not continuous

→ Also makes sense, because we did not care

Problematic for generative models.

→ What happens **here**? Never seen for decoding!

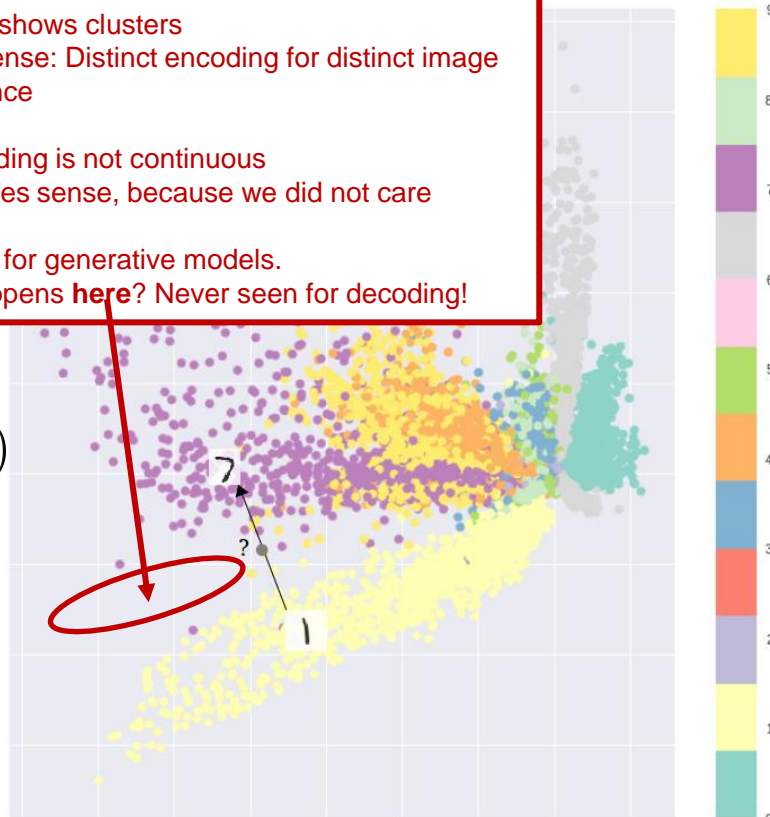


Image from [this blog post](#).

Variational Autoencoders: A Probabilistic Spin

Assume training data $\{\mathbf{x}^{(i)}\}$, $i = 1, \dots, N$ is generated from latent representation \mathbf{z}

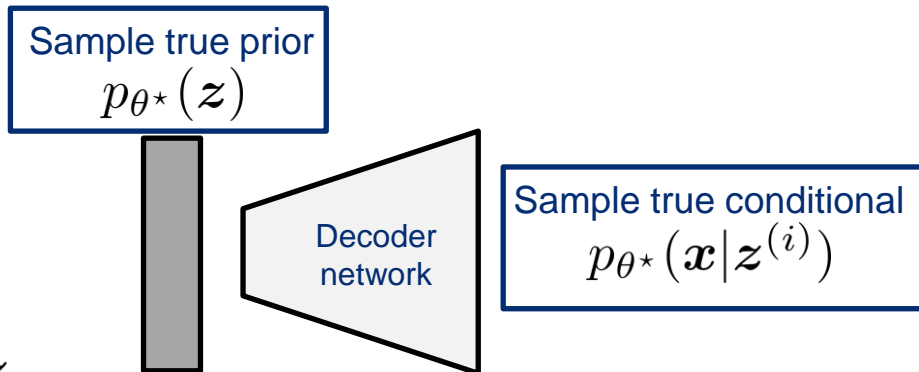
Estimate parameters: θ^*

Q: How to train this model?

Learn parameters that maximize likelihood of training data!

$$p_{\theta}(\mathbf{x}) = \int p_{\theta}(\mathbf{z}) p_{\theta}(\mathbf{x}|\mathbf{z}) d\mathbf{z}$$

Q: Any problem with this?



Variational Autoencoders

Deriving the log data likelihood

$$\begin{aligned}\log p_{\theta}(\mathbf{x}^{(i)}) &= \mathbf{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})} \left[\log p_{\theta}(\mathbf{x}^{(i)}) \right] \quad (p_{\theta}(\mathbf{x}^{(i)}) \text{ does not depend on } \mathbf{z}) \\&= \mathbf{E}_{\mathbf{z}} \left[\log \frac{p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}) p_{\theta}(\mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x}^{(i)})} \right] \quad (\text{Bayes' rule}) \\&= \mathbf{E}_{\mathbf{z}} \left[\log \frac{p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}) p_{\theta}(\mathbf{z})}{p_{\theta}(\mathbf{z}|\mathbf{x}^{(i)})} \frac{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})}{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})} \right] \quad (\text{Multiply with constant}) \\&= \mathbf{E}_{\mathbf{z}} \left[\log p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}) \right] - \mathbf{E}_{\mathbf{z}} \left[\log \frac{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})}{p_{\theta}(\mathbf{z})} \right] + \mathbf{E}_{\mathbf{z}} \left[\log \frac{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})}{p_{\theta}(\mathbf{z}|\mathbf{x}^{(i)})} \right] \quad (\text{Logarithms}) \\&= \mathbf{E}_{\mathbf{z}} \left[\log p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}) \right] - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)}), p_{\theta}(\mathbf{z})) + D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)}), p_{\theta}(\mathbf{z}|\mathbf{x}^{(i)}))\end{aligned}$$

Decoder network, estimates this term through sampling. (Differentiable via re-parametrization, see paper)

KL between encoder and z-prior →
Two Gaussians, closed-form solution

$p(\mathbf{z}|\mathbf{x})$ intractable, this cannot be computed. But $\text{KL} \geq 0$

Embedding






Image from [this blog post](#).

Unsupervised and Representation Learning

Re-introducing the split in feature extraction and semantic mapping, but trained end-to-end.

$$\min_{\theta, W} \frac{1}{N} \sum_i L(g_W(f_\theta(x_i)), y_i)$$

→ Learn generalizable, semantic representations of samples solely from data.

- Input x_i with label y_i (**Caveat:** y is derived from data itself)
- f_θ computes high-level representations of x_i
- g_W is a classifier that maps representations to labels

→ **Task: Find parameters Θ, W such that loss is optimal**



Adrianna
@adri_holmes00

My professor teaching to class on zoom:



Me, trying to prevent my professor from teaching into the void:



Adrianna
@adri_holmes00

My model, ready to just learn the identity mapping:




Me, trying to get my autoencoder to learn meaningful representations



When your GAN suffers from mode collapse



Generative Adversarial Networks




$$\min_G \max_D V(D, G) = \mathbf{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbf{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

The generator

- Accepts an input noise vector with prior $p_{\mathbf{z}}(\mathbf{z})$
- Represents mapping $G_{\theta_g}(\mathbf{z})$ that generates distribution p_g over the data space
- G is a differentiable function

Generative Adversarial Networks




$$\min_G \max_D V(D, G) = \mathbf{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbf{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

The discriminator

- Accepts an input sample \mathbf{x}
- Represents mapping $D_{\theta_d}(\mathbf{x})$ that yields probability of $\mathbf{x} \sim p_{\text{data}}$
- D is also a differentiable function

Generative Adversarial Networks




$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

The value function

- Discriminator $D_{\theta_d}(\mathbf{x})$ assigns correct label to any sample it is presented: **real** / **fake**
→ Should be maximal: Very good discrimination

Generative Adversarial Networks




$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

The value function

- Discriminator $D_{\theta_d}(\mathbf{x})$ assigns correct label to any sample it is presented: real / fake
→ Should be maximal: Very good discrimination
- Generator $G_{\theta_g}(\mathbf{z})$ attempts to “fool” the discriminator
→ Should be minimal: Poor discrimination

[Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., ... & Bengio, Y. \(2014\). Generative adversarial nets. NeurIPS \(pp. 2672-2680\).](#)



- Near convergence: $p_g(G)$ is similar to p_{data} , and $D(x)$ is partially accurate
- Inner loop: $D(x)$ is trained to better discriminate, converging to $D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_g(x)}$
- After G update: Gradient of $D(x)$ has guided $G(z)$ to be more likely classified as “real”
- After multiple iterations, $p_g(G) = p_{\text{data}}$ and $D(x) = \frac{1}{2}$

[Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., ... & Bengio, Y. \(2014\). Generative adversarial nets. NeurIPS \(pp. 2672-2680\).](#)

Challenges

Poor discriminator learning schedule

Poor generator learning schedule

Challenges

Poor discriminator learning schedule

- Discriminator provides gradients for learning
- If discriminator too strong → Saturation, and zero gradient
 - This problem is very real!
 - Beginning: Generator is weak; consider corruption etc. to confuse discriminator

Poor generator learning schedule

Challenges

Poor discriminator learning schedule

- Discriminator provides gradients for learning
- If discriminator too strong → Saturation, and zero gradient
 - This problem is very real!
 - Beginning: Generator is weak; consider corruption etc. to confuse discriminator

Poor generator learning schedule

- Generator updates w.r.t. discriminator gradient
- Will exploit any discriminator weakness → Mode collapse
 - This problem, unfortunately, is equally real
 - NNs are surprisingly brittle (will see later)

The bottom line: GANs are nice – after the fact.